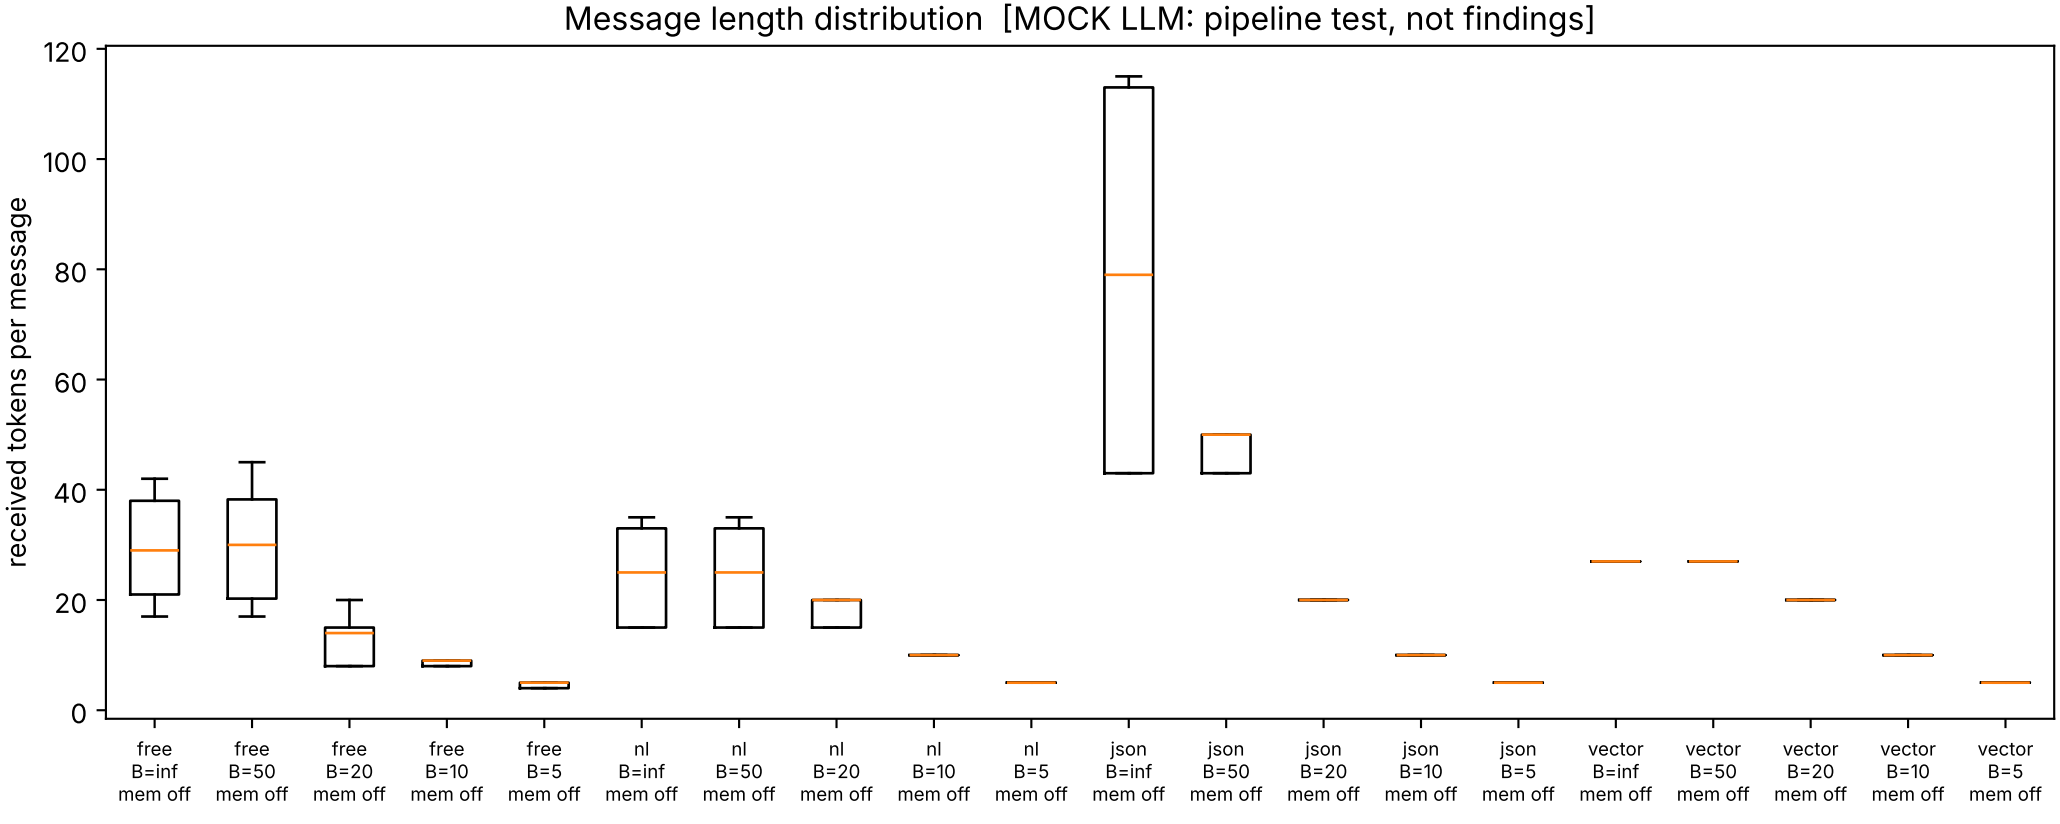
Message length distribution [MOCK LLM: pipeline test, not findings]
received tokens per message
0
20
40
60
80
100
120
free B=inf mem off
free B=50 mem off
free B=20 mem off
free B=10 mem off
free B=5 mem off
nl B=inf mem off
nl B=50 mem off
nl B=20 mem off
nl B=10 mem off
nl B=5 mem off
json B=inf mem off
json B=50 mem off
json B=20 mem off
json B=10 mem off
json B=5 mem off
vector B=inf mem off
vector B=50 mem off
vector B=20 mem off
vector B=10 mem off
vector B=5 mem off